


WEEK 1

What is Reinforcement Learning?

- **Agent:** The learner or decision-maker.
- **Environment:** Everything the agent interacts with.
- **State:** A specific situation in which the agent finds itself.
- **Action:** All possible moves the agent can make.
- **Reward:** Feedback from the environment based on the action taken.



How Reinforcement Learning Works?


- RL operates on the principle of learning optimal behavior through trial and error.
- The agent takes actions within the environment
- Receives rewards or penalties
- Adjusts its behavior to maximize the cumulative reward.

How Reinforcement Learning Works?

- **Policy:** A strategy used by the agent to determine the next action based on the current state.
- **Reward Function:** A function that provides a scalar feedback signal based on the state and action.
- **Value Function:** A function that estimates the expected cumulative reward from a given state.
- **Model of the Environment:** A representation of the environment that helps in planning by predicting future states and rewards.

Example

- **Environment:** It can be anything such as a room, maze, football ground, etc.
- **Agent:** An intelligent agent such as AI robot.



Lorem ipsum

S_6 block, which is a **wall**, S_8 a **fire pit**, and S_4 a **diamond block**.

agent reaches the S_4 block, then get the **+1 reward**

reaches the fire pit, then gets **-1 reward point**

four actions: **move up, move down, move left, and move right.**

any path to reach to the final point - fewer steps.

OpenAI gym Library

Introduction

- Open-source toolkit in Python
- Robust and scalable environment for creating, simulating, and solving RL problems.
- Key Features
 - Wide Range of Environments
 - Standardized API
 - Custom Environment Creation
 - Active Maintenance
 - Integration with RL Libraries

Installation

pip install gymnasium

Basic Usage

```
import gymnasium as gym

# Create an environment
env = gym.make("CartPole-v1")

# Reset the environment
state, info = env.reset()

done = False
while not done:
    # Take a random action
    action = env.action_space.sample()
    state, reward, done, truncated, info = env.step(action)

env.close()
```


Environments

- **CartPole-v1**: Balance a pole on a moving cart.
- **MountainCar-v0**: Drive a car up a steep hill.
- **Pendulum-v1**: Swing a pendulum upright
- **Taxi-v3**: Pick up and drop off passengers at the right locations.
- **FrozenLake-v1**: Navigate a grid while avoiding holes.
- **Blackjack-v1**: Play a game of blackjack.

Basic concepts

Environment An environment is where the agent interacts to learn behavior. It encapsulates:

- **State (Observation):** The current situation of the environment.
- **Action:** What the agent can do in the environment.
- **Reward:** Feedback received from the environment after taking an action.
- **Episode:** A sequence of states, actions, and rewards, starting from an initial state and ending in a terminal state.

```
import gymnasium as gym
env = gym.make("CartPole-v1")
```

Action Space

Defines the set of all possible actions an agent can take. It can be:

- **Discrete:** A finite set of actions (e.g., move left or right).
- **Continuous:** Actions with continuous values (e.g., steering an angle).

```
print(env.action_space)
```

Observation Space

Defines the type, structure, and range of observations (or states) the environment provides. It can include:

- Vectors of numbers (e.g., position, velocity).
- Images (e.g., frames in Atari games).

```
print(env.observation_space)
```

Step Function

The `step(action)` method advances the environment by one timestep based on the action taken by the agent. It returns:

- **State (Observation):** The new state after the action.
- **Reward:** The immediate reward for the action.
- **Done:** A boolean indicating if the episode is finished.
- **Truncated:** A boolean indicating if the episode was terminated due to a time limit or custom condition.
- **Info:** Additional debugging information.

```
action = env.action_space.sample()
```

```
state, reward, done, truncated, info = env.step(action)
```

Reset Function

Resets the environment to its initial state, starting a new episode.

Returns:

- **Initial State:** The starting observation.
- **Info:** Metadata about the environment state.

```
state, info = env.reset()
```

Reward Function

The reward function provides feedback to the agent for its actions. It guides the agent toward achieving its goal. Rewards can be:

- **Sparse:** Provided only when the goal is achieved.
- **Dense:** Provided for every step, proportional to the progress.

Episodes

An episode begins with `reset()` and ends when:

- The done flag is True (goal achieved or failure).
- The truncated flag is True (time limit or predefined conditions).

Rendering

Allows visualization of the environment to see what the agent "sees."

```
env.render()
```

Closing an Environment

Always close the environment after use to free up resources.

```
env.close()
```

Steps

1. Definition:

A step refers to a single interaction between the agent and the environment. During a step:

- The agent takes an action.
- The environment responds with the next state, a reward, and information about whether the episode has ended.

2. Key Characteristics:

- Represents a **single timestep** in the environment.
- Involves the transition from one state to another.
- Cumulative reward and learning updates occur at each step.
- Steps accumulate within an episode until it terminates.

3. Example: For a CartPole environment:

- At step 1: The agent moves the cart to the right.
- At step 2: The agent moves the cart to the left.

Episode

1. Definition:

An episode is a complete sequence of steps, starting from the initial state (reset) and ending when the environment reaches a terminal or truncated state.

- **Terminal:** Achieving the goal or failing (e.g., the pole falls in CartPole).
- **Truncated:** Reaching a maximum number of steps or custom-defined criteria.

2. Key Characteristics:

- Represents a **full trial** or run of the environment.
- Ends with either success, failure, or a predefined condition.
- Multiple episodes are required to train an agent, as learning happens over time.

3. Example: For the CartPole environment:

- An episode might start with the pole upright.
- The agent interacts with the environment, balancing the pole through multiple steps.
- The episode ends when the pole falls or a time limit is reached.

WEEK 1: TUTORIAL ON OPENAI GYMNASIUM

1. Take the tutorial: Getting Started With OpenAI Gym: The Basic Building Blocks
 - <https://www.gymlibrary.dev/content/tutorials/>
 - <https://blog.paperspace.com/getting-started-with-openai-gym/>
2. Use the CartPole-v0 environment and write a program to :
 - a. Implement the CartPole environment for a certain number of steps
 - b. Implement the CartPole environment for a certain number of episodes
 - c. Compare and comment on the rewards earned for both approaches.
 - d. Plot the cumulative reward of the games and note down the results.

Implement the same considering the
“Mountain Car” environment